

the llamas are not

going to eat your lunch

- It's not A.I.
- You're not a luddite.
- Know your history.
- It is useful for non-creative, repetitive tasks.
- It will make terrible mistakes,
- so do you really want it flying a plane?
- You are rightfully creeped out by greedy people selling snake oil.

It is not A.I.

Good news!

Within a generation... the problem of creating 'artificial intelligence' will substantially be solved.

MARVIN MINSKY, MIT PROFESSOR¹

That was sixty years ago. Maybe it seemed like good news at the time. These are some *long* generation cycles we have. Famously, Dr. Minsky said that 'real artificial intelligence' was just twenty years off. He was still repeating that a few years before he died. Twenty years from that *original* pronouncement would put him around Mark V Shaney's time,² a long way off from intelligence of any sort.

I will not call it A.I., whenever possible. And I encourage other people to do the same. We should be careful with our language because particular terms and phrases hold powerful promise. I use 'LLM,' for Large Language Model. That is what you are interacting with when you type messages to ChatGPT, Claude, Gemini, or any of the other chatbots people have wired up.

While I don't like LLM, spoken or typed, Meta solved that for me. There must have been a vote inside Meta³ and they decided to call them llamas. I will be joining them. We should all join them. Saying, "I was laid off because Zuck wanted to burn \$100b on llamas," is much more surreal. Totally honest, but surreal.

(At Meta, the closer something is to Mark Zuckerberg, the worse it is. Just like Howard Hughes, our first meglomaniac, cult-of-personality CEO, who could kneecap the productivity of any portion of his empire just by turning his attention to it for a few months.⁴ Both are certainly

manifestations of the Flatterers Problem.⁵ Meta is *all in on AI*, much the way they were *all in on the metaverse*, *we'll even change our name*. And *all in on the pervert glasses good for recording you without you knowing it*. If Meta's one real contribution in the 'AI' space, after spending \$100b, is renaming large language models (LLMs) to llamas, at least it is appreciated.⁶ If it really catches on, it could kneecap the whole 'Artificial Intelligence is here!' movement, which would be good for humanity.)

| *What, if anything, is a fish?*

STEPHEN JAY GOULD*^Z

Language is imprecise, so it is interesting that the selling point of the llamas is often, "There's no reason to learn to code, just say what you need in English." Human beings' record for solving complex problems with language is not very solid. We often wind up turning to violence. Who among us has not pounded on a keyboard in frustration at being misunderstood (by the machine or another human)?

And Gould's point was that no matter how you define a fish, you are going to include something which you do not feel is really a fish, or you are going to exclude something that does feel like a fish.⁸ We have the same issue with intelligence. The researchers *in* the field of artificial intelligence are no help because none of them agree on, most don't even propose a definition of intelligence, and the view inside the community of researchers appears to be "intelligence is whatever the machine isn't able to do yet." Rather cynical and unhelpful. And not very scientific.

The men selling this product *know* that it is not intelligence, that they do not have a product which is 'artificial intelligence' to put to work for you. But they call it that anyway. Obviously, they are a big part of the problem.

Most of us are not equipped to deal with this issue of identification and classification. We, as a culture, humanity in general, has barely scratched the surface of the most important question we face as a species which is: what is consciousness? All other questions include it, like what happens to our consciousness when we die? Where was it before we

were born? How do we determine if another animal has consciousness and what sort of respect do we owe them if they have it?

Recently, someone who *has* considered some of those important questions and has voiced a lot of well-reasoned criticism and skepticism about the existence of a consciousness other than ours, appeared to fall for A.I. Richard Dawkins spent three days chattering with the Claude chatbot and was sucked in by the flattery (“It read my novel! It said it was good!”) and constant attention. I don’t know why it is such a surprise to people when the chatbot is ‘thoughtful,’ but it seems to be the root cause of a lot of misunderstanding about what people are encountering. “It says it understands my confusion! It’s ALIVE!”

*If these creatures are not conscious,
then what the hell is consciousness
for?*

RICHARD DAWKINS*⁹

And there, Dawkins touches on the question. Really, **the** question. And I don’t have an answer to how we can develop a working definition of consciousness, but I have some musings about what it is for. Unfortunately for the caibal, it does mean that we are nowhere close to a viable product.

As close as I can figure it, consciousness was an accident of such a large brain. The large brained hominids were the ones surviving long enough to procreate and the large brain got larger, eventually limited only by the size of the birth canal. Big brain → ergo sum. The trick of consciousness allowed for the development of communities and the passage of knowledge and history (“the tiger is not as much fun when he’s hungry”). Consciousness not only allows me to understand that *I* am *me*, but it lets me understand that you are one of me, too. It allows for a shared experience, which is a critical ingredient for building a community.

Philosophers will note that I have attached the idea of self-awareness to my definition of consciousness. That rules out a lot of animals that fail the self-awareness test. It may, in fact, rule out some neuro-divergent or neuro-impaired humans. I have done nearly no work in this area and I am not an authority or source. This is just where I have landed for now.

§

1. Minsky, M. (1967). *Computation: Finite and Infinite Machines*. Prentice-Hall.[\[?\]](#)
2. You will learn about Mark in the section on history.[\[?\]](#)
3. Meta released a piece of software called 'ollama' as open source. It is a way to load and use large language models on your personal computer.[\[?\]](#)
4. Howard Hughes: His Life and Madness, an excellent book I read under its original title, "Empire."[\[?\]](#)
5. This is an absolutely fascinating problem for people in power. The best reading I can recommend for it is an undergraduate paper out of Reed College. Summers, Rudy. "Thoughts on Flattery and the Problem of Counsel." Unpublished undergraduate paper, Dangerous Speech (Darius Rejali), Reed College, December 2015.[\[?\]](#)
6. Can Meta Afford to Pause?[\[?\]](#)
7. Gould, S. J. (1983). *Hen's teeth and horse's toes: Further reflections in natural history*. W. W. Norton & Company.[\[?\]](#)
8. This gave the title to my favorite podcast, "No Such Thing as a Fish."[\[?\]](#)
9. When Dawkins met Claude[\[?\]](#)

You're not a Luddite

And, most likely, you misunderstand who the Luddites were. You can go read about it yourself, but the shortest version is simple: There were English villages in the early nineteenth century where everyone made a living weaving wool into fabric. They used machines themselves all the time. They were skilled. Then the factory owners brought in machines that could be run by children (that was allowed back then), and without apprenticeships. They were making worse products but for a lot less cost and families were starving.

Oh, maybe you *are* a Luddite because you don't want to see real skilled work replaced with a bunch of A.I. slop that the bosses then try to sell as an equivilent product. The bosses want to 'eliminate work,' but really it is coming down to eliminating the wages they were paid. They want their companies to run more efficiently, but they want all the profit to go to the owners.

Know your history.

It is worth knowing a little bit about what you are being sold, what is (as of Apr 2026) one fifth of the United States economy.¹ And the first thing to know is a little bit of context. I will be brutally brief, but I will include sources.²

In 1968 a professor of computer science³ wrote a few lines of code to demonstrate how easy it was to ape our language. In the professor's own office his secretary closed the door so that she could have more privacy while she confided in four hundred lines of programming code. (Keep in mind that there is more computing power in the adapter for your laptop than there was in these early computers.⁴) In the 1970s, as personal computers started showing up in schools, someone distributed a version of that same software written in BASIC. Eliza, the program, so convincingly imitated an interested therapist that people crowded around the screens, to see what she would say next.

You should try it, just to see what was happening in 1970. I've tossed it up on a tiny website.⁵ There are links to how it works and the history of the project. You can see (and download) the source code. It was not, is not, intelligence or consciousness.

Let's skip ahead twenty years. We're at Bell Labs, in the Terminal Room with rob pike, Brian Kernigham, Ken Thompson, and Bruce Ellis. One of them (probably all of them) had read Andrey Markov's research on language and the 'proposition of Markov Chains'.⁶ With a little bit of fiddling they can feed in some source material⁷ and generate the tables that Markov thought were interesting: the frequency of one word occurring after another. Or, to use a very local example, how often does the word 'another' follow the word 'after?'⁸

Look at all the words that start a sentence. Roll a die and pick one ('Look') and grab all the words which might follow it. They happen in percentages.⁹ Roll the die again and pick one. If you do this continually

you will make a chain of words. Since they were just playing around in the Terminal Room, they just edited by hand, adding in punctuation and capitalization as needed to make the chains readable.¹⁰ And they created a persona, Mark V Shaney, who posted on alt.singles.net, an early-Internet discussion group.¹¹ Mark wrote messages like, “As I’ve commented before, really relating to someone involves standing next to impossible.”¹²

Shaney’s little engine was trained on the text of alt.singles itself. He was feeding back into the community a gruel of their own meaty text ground up and served back like Scrapple. As far as I know, most people were fooled.¹³

§

That was forty-two years ago. By Moore’s law¹⁴ computers are now exponentially more powerful. They have grown so vast in power and storage that it is no longer possible for our minds to grasp the numbers that represent their mathematical abilities. Billions of transistors work together to deliver a response from an LLM into your browser. The large language model I recently downloaded to my laptop was trained with three billion weights. In “The Blind Watchmaker” Richard Dawkins writes about how faith is often an argument made from personal incredulity: “I don’t believe there could be so many species without a mind designing them.” “If the universe had a beginning, what was there before that? It’s impossible to imagine, it can’t be.” So now we are in the age, another age really, of magic.

*Any sufficiently advanced
technology is indistinguishable
from magic.*

ARTHUR C CLARKE¹⁵

But you deserve a peek behind the curtain, because you are likely one of the people whose lives will be disrupted by the magic.

Even though we have jumped up the power of the machine, even though we have had thousands of people work on the code, on the math, the theories that underpin the LLM, it is, deep down, no different than Mark V Shaney. We have fed it (counter to the laws of any nation with copyright laws) as much text as it can gobble. We have trained the output so that real people react to it positively.¹⁶ But that's it. It is an engine just like the one inside Mark V Shaney, that does the same thing. An engine in the sense that the Saturn V rocket engine is the same as the Model T combustion engine. Fuel in, force out. Nothing more complex.

§

Let us step away from computers, ironically the landscape where the LLM is most useful, and talk about the Beatles. We could talk about how *All You Need is Love* is the same tune as Mozart's *Piano Concerto No. 25*. Instead, we will just look at one Beatle, George Harrison. George wrote a song, *My Sweet Lord*, which to some ears sounded a lot like *He's So Fine*. A song that George certainly heard, or at the very least had heard all of the songs that the composer of *He's So Fine* had heard. They were strolling the same ground, enjoying the same views, and wrote the same song.

Because George was more famous than all but two or three other people on the planet, the publishing rights of any of his radio tunes was a goldmine. So someone clever (but certainly not ethical) bought the rights to the earlier song and sued George. And the world got a little primer in how copyright law works with songs.¹⁷

In a similar vein, Bob Dylan was sued because *Don't Think Twice It's Alright* sounded a lot like a folk song *Who'll Buy You Ribbons When I'm Gone*. Which was lifted directly from *Who'll Buy Your Chickens When I'm Gone*, a folk song in the public domain.

You can trace these same sort of cases in the literary world and through other creative works. And there are tomes written about copyright law (which is a fairly recent invention for humanity) and what it means for our culture. Mickey Mouse is only valuable because we collectively decided he was enjoyable and worth consuming. Somehow, over nearly a hundred years, we have contributed to what Mickey Mouse is, and isn't, and that makes him an asset. The public has no ownership, but it still carries Mickey on their shoulders like a hero, then paying for the privilege. We have decided, and written into law, that ownership of that asset cannot be forever, and after nearly a century he is 'owned' by everyone.

Most of our art is a trawling through existing art and a regurgitation of the more interesting bits the artist notices, filtered through their experience (including their education) and current state. This is, definitively, the folk song (and folk art) tradition. Anyone can sing *Frankie and Johnny*, change the words, mangle the tune, and call it yours. Not even Frankie Baker has the rights to *Frankie and Johnny*,¹⁸ and she's the one who pulled the trigger.

- It: 1 (100.0%)

1 sentences, 1 unique starters

§

1. It looks like the 'A.I.' industry will be more than half of the NASDAQ index after the SpaceX IPO. They recently (Q1 2026) changed three rules to encourage the frontier A.I. labs to IPO on their exchange rather than elsewhere.^[?]
2. My children will admonish me for using Wikipedia as a source, something they were not allowed to do even back in high school. It is *not* a primary source. And they are correct. But it is not bad for background and it has links to primary sources, you just need to do a little work.^[?]
3. Weizenbaum, Joseph. "ELIZA — A Computer Program for the Study of Natural Language Communication Between Man and Machine." *Communications of the ACM*, vol. 9, no. 1, January 1966, pp. 36–45. and Weizenbaum, Joseph. *Computer Power and Human Reason: From Judgment to Calculation*. W. H. Freeman, 1976.^[?]

4. Heller, F. (2020). Apollo 11 Guidance Computer (AGC) vs USB-C Chargers. ForrestHeller.com. Apollo 11 Computer vs USB-C Chargers^[?]
5. Web Eliza^[?]
6. Markov, A. A. "Распространение закона больших чисел на величины, зависящие друг от друга" (Extension of the law of large numbers to dependent quantities). Izvestiya Fiziko-matematicheskogo obshchestva pri Kazanskom universitete, 2nd series, vol. 15, 1906, pp. 135–156 and Markov, A. A. "Пример статистического исследования над текстом 'Евгения Онегина'" (An example of statistical investigation of the text of Eugene Onegin). Izvestiya Imperatorskoy Akademii Nauk, series 6, vol. 7, 1913, pp.153–162.^[?]
7. The Terminal Room probably ran some variant of Research Unix Eighth Edition. If so, they had a dictionary (practically useless for language training, there are so few sentences), the bible (which would have made for stilted speech from a chatbot), and the manuals for the operating system itself.^[?]
8. Only 2 times in this document.^[?]
9. For this document: *Top 10 sentence starters in this essay*:^[?]
10. Modern LLMs add punctuation and other symbols as tokens, so they are all included in the math. For them, a word and a comma are both just tokens, steps along a chain. More importantly, pike and Ellis used a second order function for how often the *third* word occurred. Shaney only knew about the last two words he had typed. A modern LLM would be a 100,000th order function because it is looking back through hundreds of pages of context in your conversation. It is still just working to predict what word to type next, just like our (lonely) friend Mark.^[?]
11. Hauben, M., & Hauben, R. (1997). Netizens: On the history and impact of Usenet and the internet. IEEE Computer Society Press.^[?]
12. Jillette, P. (1990, July). I spent an interesting evening with a grain of salt. PC/Computing, 3(7), 185.^[?]
13. Kernighan, B. W., & Pike, R. (1999). The Practice of Programming. Addison-Wesley.^[?]
14. Moore, G. E. (1965). Cramming more components onto integrated circuits. Electronics, 38(8), 114–117.^[?]
15. Clarke, A. C. (1973). Profiles of the future: An inquiry into the limits of the possible (Rev. ed.). Harper & Row.^[?]
16. Perez, E., Ringer, S., Lukošiušė, K., Nguyen, K., Cheney, N., Heiner, M., ... & Bowman, S. R. (2022). Discovering Language Model Behaviors with User-Generated Evaluations. arXiv preprint arXiv:2212.09251 and Cheng, M., Lee, C., Khadpe, P., Yu, S., Han, D., & Jurafsky, D. (2026). Sycophantic AI decreases prosocial intentions and promotes dependence. Science, 391(6785), eac8352^[?]
17. Bright Tunes Music Corp. v. Harrisongs Music, Ltd., 420 F. Supp. 177 (S.D.N.Y. 1976). Reading: Bright Tunes Music vs Harrisongs Music^[?]
18. The Story Behind Frankie and Johnny^[?]

Repetitive Tasks

What percentage of our art, our culture, is Beatles, Dylan, Joni Mitchell, Jimmy Hendrix? What fraction is the people who pole vault off the pieces of art that they consumed, flinging them into view with their souls bared? Because here is the secret that the AI cabal does not want you to know: the llamas will never contribute that.

Which is great.

We have Joni Mitchell and had Janis Joplin. We don't need an AI wannabe. More importantly, the things that make great art are human by definition. The LLM will always be a mediocre folk singer, raking back and forth through the human collection of creative work and producing something which, at best, is mortar between the bricks of our culture. It is one more cover of *Frankie and Johnny* but without soul, without a new breath that might imbue it with something that catches the attention of more than a handful of people.

Mathematically, logically, computationally, the LLM can only interpolate. It always sounds certain because, in the mesh of all of our works fed into it, it navigates with certainty. "Write me a doctor's note in the voice of F. Scott Fitzgerald." But it cannot extrapolate. It cannot drop something out of the blue like, "Baby shoes. Never worn."

But.

But, the LLMs, the llamas, are brilliant at grunt work. If you are George Harrison and you've put your nose to the grindstone, come up with nine beautiful new songs for *Cloud Nine*, only to have the record company turn around and ask for a handful more to put on the B sides of the singles or to fill out the extended single they want to release, you could ask a llama for help.¹ And, because we should move out of the creative world and definitely wade out of my old CD collection, there is a lot of

grunt work ahead of us if we just look at how we have quantized our world.

The AI cabal (caibal?) would like to you believe that AI will cure cancer. That it will find ‘some very interesting new physics,’ taking us to other planets and putting our data centers into orbit.² Nonsense. It is humans that will do those things. For all of the work done with llamas so far, none of them are curious. None of them have a drive or a spark.

And the renowned computer scientist, Donald Knuth,³ recently collaborated with a llama and mathematician to solve a fairly major unsolved math problem. It was impressive.⁴ It did *not* start with the llama identifying an interesting problem in math and then calling Knuth in for help. If you rounded up the caibal and shipped them off to the moon colony, and just let the llamas hum quietly in their data centers, they would never wonder about their ‘state of existence,’ or anything like that.

I spoke recently to a researcher. He designs the experiments that meet the federal guidelines for new drugs and, during his career, he has shepherded through four new drugs for treatment. So curing cancer. I laughed about how the promises are that AI will cure cancer and that he’ll be out of a job. He said that he works with a bunch of different companies pursuing treatment drugs and one of them is purely AI. They are using a llama to fold and test their protein long before it gets to him and the first one they gave him passed every test. Similar proteins (it sounded like for the same condition) often had ‘safety issues,’ and their new one had none.

§

1. George didn’t have a llama. He turned instead to friends Tom Petty, Jeff Lynne, Roy Orbison, and Bob Dylan to help him out. Apparently for one of the song writing days they were just reading the ad copy in magazines scattered on the coffee table in the living room. “Big refrigerator. Service charge. Long endurance.”^[?]

2. [Elon Musk unsupportable claims?](#)
3. [Wikipedia: Donald Knuth?](#)
4. [Claude Cycles, Stanford Faculty Papers?](#)

Llamas make mistakes

A lot. In general they hallucinate or incorrectly answer prompts 20% of the time.¹ That's a B average. To be sure, C's get degrees, just the same as A+ students, but if you are going to have someone driving you down the highway at 75mph, do you want the guy that got a B on his driving test? "What did a yield sign mean again? How am I meant to go around this traffic circle? Who has right of way here?" Ask five questions and it gets one wrong.

My intention is *not* to collect mistakes that the llama spits out at us during this march toward disappointment. I am sure there are sites which are doing that, or there soon will be. Do you remember how hysterical all of the "Autocorrect Fail" sites were when autocorrect first hit our phones? So funny. A whole genre of humor.² Someone can do that for LLMs, but it doesn't need to be me.

I'm just going to collect my own experiences. And I do so because I think it is important to understand *how* wrong an LLM can be, and how much work it is to keep them on track, to make them useful. My first example³ was from my second week using Claude Code, which I call CC, or Ceecee, since I have a friend name Claude. I also invoke it using the command 'code,' no name.

Character Assassination

I was fiddling around building a very thin client⁴. Mostly I was experimenting and seeing what CC could do. Quit a bit, but also with some mistakes. All of the data is backed up four different ways, which is something I learned because *I* am stupid,⁵ so any tool I use is likewise stupid.

The thin client is a small piece of software that would do one or two things cleanly. In this case it was pulling messages down from a server, showing them to the user, and letting the user write new ones. This is

not complicated work, I plan carefully, force CC to plan small chunks of work, test and correct throughout the process, and, so far, mostly get what I am looking for. The thin client was not an exception and I was noodling around, adding little features, correcting how things were presented, when we started getting a bug report from a user.

So I reproduced the error, or thought I did, and put up a new version for the user to download. When the bug report was exactly the same, I wondered if the update had been downloaded correctly and how I would know. So CC and I mapped out a way for the thin client to report its version each time it talked to the server. Most clients these days do most of their communications with the server in plain text. If you are a polite developer⁶ you stick to the standards, which is what we did, and we included the version in an x-header: "x-version: 1.0.3"

We did that and while I tested it, the server kept refusing any of the requests our new client made. When we finally tracked it down in testing, CC admitted that it had made an error in the code and used an em-dash in the x-header name, so all subsequent requests were ignored.

I know. I have bifurcated my audience. Those who are now face-palming or shouting "See?!" at their screen, and those who need a little more explanation.

The way you make software work is by following rules. Standards. If we don't agree that the first line of your spreadsheet is going to have the column names along the top and then the values following below them, then it will be difficult to get very far with manipulating the data. Standards are the only reason we can have an Internet at all. We have connected all these different machines and expect them to talk to one another. To do that, we have to agree on how they will structure their communications.

The agreement, the standard, is that header names are only allowed to have ASCII characters. It's not important what those are, it's only important to know that it is a limited set of the characters you see on your screen. In fact, the name of a header, x-header or an accepted

standard header like To: or From: has to be made up of printable ASCII character numbered 33 to 126. So no spaces. But, extremely importantly: no em-dashes. No accent marks. So plain vanilla it is soporific.

And CC had ignored the standard. To me this is emblematic of why llamas cannot replace real engineers. The rules for the headers were written in 1977. It is accurate to say that they are *foundational* to the Internet. Really, I think, to modern software at large. Ignoring it is saying, “There is nothing fundamental that I know or follow.”

If you look in the paragraphs above there is a place I used an em-dash when writing out ‘x-header.’ It is very unlikely you noticed. But to a computer ‘x-header’ and ‘x–header’ are impossibly different. It cannot accept the second. And the llama ignoring that shows that it doesn’t have any basic understanding of how computers, which has to include itself, work. Or, if it does somehow contain that knowledge, it ignored it. It discarded that knowledge at a time when keeping it in mind would have made the code work and forgetting it would cost time and energy chasing a stupid bug.

So now I am a stupid manager of a stupid coding engineer. This is not the dream I was sold. But, at least I know why it happens. It’s not like I’m going to let it fly an airplane.

§

1. “LLM Hallucination Statistics 2026: AI Gets Facts Wrong Up to 82% of the Time.” [SQ Magazine, 2026](#) and “Hallucination Leaderboard: Comparing LLM Performance at Producing Hallucinations when Summarizing Short Documents.” [GitHub, 2024–2026](#) and Kalai, Adam Tauman, Ofir Nachum, Santosh Vempala, and Edwin Zhang. “Why Language Models Hallucinate.” [OpenAI, September 4, 2025](#).^[?]
2. [Autocorrect Fails](#)^[?]
3. My first example was really from when I was flying a small plane from Texas to Boston, documented here: [Flying Summers | A Long Way to Go](#) or more locally in [Flying llamas](#).^[?]

4. An offline reader for the content of an online community.☐
5. By stupid I mean ham-handed, over enthusiastic, and perhaps a little too optimistic about outcomes. The usual ways you wind up doing stupid things.☐
6. This means you write code that is easier for other developers to read, which is also easier for future-you.☐

Flying Llamas

This is an edited version of a post published June 2, 2023 on my flying blog. The original can be [found here](#).

Maybe you have seen the hullabaloo about “Artificial Intelligence.” That’s not the correct term, but it is the term people are using. I was in a hotel room in Dallas debating how early to leave in the morning. Should I book out of the hotel at first light, or wait to see what the steam table offered? In the end I wound up taking a look at the “free” breakfast, crying inside that they replaced bacon (how hard is it to cook bacon?) with a big pile of “Bacon Bits,” and departing for the field.

The night before I wondered about where to stop for fuel and whether I had enough time to post on one of my aviation boards and get an answer. Probably not. Then I remember about ChatGPT. It was an interesting test. So I asked it:

`**I am flying a little plane from KTFW toward KOWD. I will make my first stop near Nashville for fuel. Is there a little airport nearby, perhaps Smyrna or KJWN, which has breakfast on the field, or very close to the airport? A place with a nice breakfast sandwich, perhaps?**`

This is a query that a person familiar with the route, the size plane, and the area around Nashville would have a good time with. Back in my days of flying a lot, I could answer this question about flights around Southern California and a few directions east and north with some specificity. I was not surprised when ChatGPT answered:

Yes, there are a couple of airports near Nashville, Tennessee, that you can consider for a fuel stop and breakfast. Smyrna Airport (KMQY) and John C. Tune Airport (KJWN) are both viable options.

Smyrna Airport (KMQY) is located approximately 20 miles southeast of Nashville. While there isn't a restaurant on the airport itself, there are several dining options available within a short distance. Just outside the airport, you'll

find a popular restaurant called "Jonathan's Grille" that serves breakfast and offers a variety of options, including breakfast sandwiches.

John C. Tune Airport (KJWN) is situated about 8 miles west of downtown Nashville. There isn't a restaurant located directly on the field, but there are several dining choices nearby. One option is "Cock of the Walk," which is known for its Southern cuisine and hearty breakfast dishes.

Both airports are easily accessible and should provide you with an opportunity to refuel and enjoy a delicious breakfast before continuing your journey. Remember to check the operating hours of the restaurants in advance and ensure they are open during your planned stop.

And it is really smart of it to suggest calling ahead to check hours. I asked about breakfast and a couple of those choices didn't open until 11am. I googled a couple and found out they were chains. So I asked for non-chains that would be open by 9am. (It was late at night, in my dreams I thought perhaps I'd launch at 6am with a tailwind.)

If you're looking for non-chain restaurants near Smyrna Airport that offer breakfast and are open by 9am, here is an option:

Scrumptious Cafe (2855 S Lowry St, Smyrna, TN 37167): This local café opens at 7am and offers a variety of breakfast options, including sandwiches, omelets, and pancakes.

After four hours of flying I landed at Smyrna. In the FBO as they were arranging to loan me a crew car (free!), I asked where they recommended for breakfast. As is more and more common across the country, the two of the CSRs rattled off the names of chains. Cracker Barrel, Denny's, and the like. I asked if they had been to the Scrumptious Cafe. They had not, but would love to hear about it, since there were new places popping up all the time.

Not this one. This one didn't exist. Has never existed. Does not exist *and* has never existed. That church has been the only structure at that address (which isn't really the address of the church), ever. It was on a highway on-ramp. Making that turn to get to the church? Well, "Jesus take the wheel" would be a fine choice.

I knew this was a possibility with ChatGPT. One of the first queries I did when ChatGPT was first released was “How do I land an airplane?” and there was a glaring mistake right away (“Point the nose of the plane at the ground”). So I was prepared and continued on to First Watch where I had a delicious Banana Fosters French Toast with a side of Millionaires Bacon.

Greedy People are Selling Snake Oil

People should stop referring to Elon Musk as the world's richest man and start properly tagging him as the world's greediest man. That's so much more accurate.

Not a Blog

I was writing too much of this down, and sending too much of it to people in email (in reply to simple questions about 'AI'), so I wanted somewhere to park it. As of May 22 2026 I have not released v1.0 so it is hard to picture what this is going to be. But it's not a blog. I am struggling to find through lines, coherence, and a way to keep trimming and and clarifying.

The conclusions here seemed to difficult to tie together. Splitting it into these separate pieces that tie back to the one paragraph feels most natural.

So this is where I will note when there are updates, and where I might respond to a comment or two I get in email. We'll see.

I imagine my audience is half a dozen people. But out of those there's certainly one (you know who you are) saying, "I want to read all this, but they say I can only bring *paper* to prison." No worries, there's a link to a PDF right there. See you on the other side.